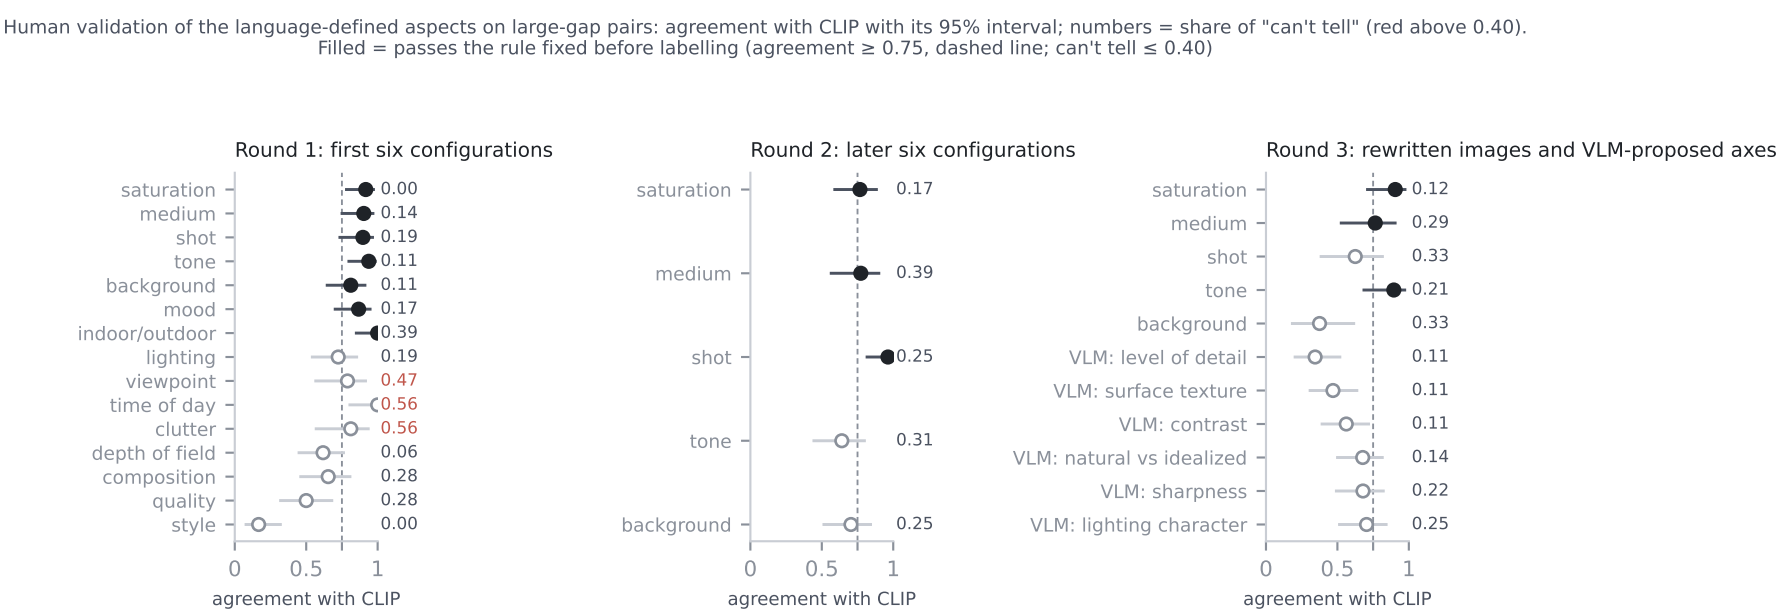
Human validation of the language-defined aspects on large-gap pairs: agreement with CLIP with its 95% interval; numbers = share of "can't tell" (red above 0.40).
Filled = passes the rule fixed before labelling (agreement ≥ 0.75, dashed line; can't tell ≤ 0.40)
Round 1: first six configurations
saturation 0.00
medium 0.14
shot 0.19
tone 0.11
background 0.11
mood 0.17
indoor/outdoor 0.39
lighting 0.19
viewpoint 0.47
time of day 0.56
clutter 0.56
depth of field 0.06
composition 0.28
quality 0.28
style 0.00
0 0.5 1
agreement with CLIP
Round 2: later six configurations
saturation 0.17
medium 0.39
shot 0.25
tone 0.31
background 0.25
0 0.5 1
agreement with CLIP
Round 3: rewritten images and VLM-proposed axes
saturation 0.12
medium 0.29
shot 0.33
tone 0.21
background 0.33
VLM: level of detail 0.11
VLM: surface texture 0.11
VLM: contrast 0.11
VLM: natural vs idealized 0.14
VLM: sharpness 0.22
VLM: lighting character 0.25
0 0.5 1
agreement with CLIP